

Sacking, revolt, return: how crisis at OpenAI over Sam Altman unfolded

Publication Date: 2023-11-25

Author: Dan Milmo

Section: Technology

Tags: OpenAI, Artificial intelligence (AI), Computing, features

Article URL: <https://www.theguardian.com/technology/2023/nov/25/how-crisis-openai-sam-altman-unfolded>



When Sam Altman, the chief executive of OpenAI, took to the stage in San Francisco nine days ago he hinted at another significant development in the world of artificial intelligence. “Four times now in the history of OpenAI, the most recent time was just in the last couple weeks, I’ve gotten to be in the room, when we sort of push the veil of ignorance back and the frontier of discovery forward, and getting to do that is the professional honour of a lifetime,” he told the Asia-Pacific Economic Cooperation (Apec) summit. Given that he leads the company behind ChatGPT – a chatbot that has transformed the debate around AI – this was a tantalising comment. And a major event in AI did occur the next day – Altman was fired. OpenAI’s board announced on Friday 17 November that it had sacked the 38-year-old for failing to be “consistently candid in his communications” with its members, without giving further details about the alleged breaches of trust. Altman was given notice of what was coming after he came off stage, receiving a text from fellow board member Ilya Sutskever, OpenAI’s chief scientist, asking the CEO to talk at noon via Google Meet the next day. Altman joined the call and was told by Sutskever that he was being fired, according to a post on X, formerly Twitter, by Greg Brockman, OpenAI’s chair and president. Brockman was not on the call but minutes later he too was summoned to Google Meet and told that he was being removed from the board. Brockman was told he could stay at the company but chose to quit. OpenAI appointed Mira Murati, the company’s chief technology officer, as interim CEO. The announcement of Altman’s sacking sent shock waves through Silicon Valley and immediately led to investors in OpenAI pushing for his return. OpenAI was founded as a non-profit company whose board oversees a commercial subsidiary in which Microsoft is the biggest investor. Microsoft’s CEO, Satya Nadella, said his company had not been contacted before the announcement. Altman posted on X that Friday that he had “loved my time at openai” but it soon became clear the story was far from over. The Information, a tech news site, quoted a memo from the company’s chief strategy officer, Jason Kwon, telling staff in a memo on Saturday that an effort was under way to bring back Altman and other colleagues who had quit, including the three senior researchers Jakub Pachocki, Aleksander Mądry and Szymon Sidor. “We are still working towards a resolution and we remain optimistic,” Kwon wrote. “By resolution, we mean bringing back Sam, Greg, Jakub, Szymon, Aleksander and other colleagues.” Sunday brought more twists. Altman posted a picture of himself inside OpenAI’s San Francisco headquarters holding a guest pass with the message “first and last time i ever wear one of these”. The talks with OpenAI’s four remaining board members failed and Emmett Shear, the former CEO of video game streaming platform Twitch, was appointed the company’s third CEO in three days. Soon after, Nadella announced that Microsoft had hired Altman and Brockman to lead a “new advanced AI research team”, although he added in a round of media interviews on Monday that he was open to both staying at OpenAI. As one analyst pointed out Altman

would have a de facto leadership role at OpenAI because of Microsoft's investment. By Monday simmering staff discontent at Altman's ousting broke into open revolt. The vast majority of OpenAI's 770 employees including Murati signed a letter addressed to the board threatening to quit en masse unless Altman and Brockman were reinstated by the board – a move that must then be followed by the board stepping down. "We are unable to work for or with people that lack competence, judgment and care for our mission and employees," said the letter, which surprisingly was also signed by board member Sutskever. The other three board members were: Adam D'Angelo, the co-founder of the question-and-answer website Quora; Tasha McCauley, a tech entrepreneur; and Helen Toner, a director and AI safety expert at Georgetown University's Center for Security and Emerging Technology. It became clear, however, that Altman was edging closer to a return. OpenAI's vice-president of global affairs, Anna Makanju, told staff that "intense discussions" were taking place to resolve the crisis. Bloomberg reported that Murati sent a message to staff on the Slack platform saying "still working on it ...". Then a deal was struck. OpenAI announced on Tuesday night in California that it had reached an agreement in principle for Altman to return with a new-look "initial" board chaired by Bret Taylor, the former co-chief executive of software company Salesforce, along with Larry Summers, the former US Treasury secretary, and D'Angelo. Altman has not, so far, regained his place on the board. OpenAI will reportedly conduct an independent investigation into the circumstances around Altman's departure, which has initially been announced by Shear. Shear had denied that Altman's sacking was related to safety concerns. But after Altman's return, Reuters and the Information reported that some OpenAI researchers had been concerned by the capabilities of a new AI model being worked on the company before the CEO was fired. According to the Information, the "Q*" model – pronounced "Q-star" – was able to solve maths problems it had not seen before, representing a potential breakthrough in the technology. OpenAI was founded in 2015 with the goal of developing "safe and beneficial artificial general intelligence [AGI] for the benefit of humanity". AGI refers to an AI that can perform a variety of tasks at, or above, a human level of intelligence. Some experts are concerned that the emergence of AGI could lead, in theory, to a system capable of evading human control and endangering humanity. OpenAI's board has changed but its mission has not. Altman, who had reportedly demanded a new board as a precondition for his return, indicated on Thursday – Thanksgiving in the US – that he was in the process of healing any divisions with D'Angelo. "Just spent a really nice few hours with @adamdangelo," he wrote on X. "Happy thanksgiving from our families to yours!" Altman remains where he was nine days ago: at the frontier of discovery.